

```

# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import pprint

from sklearn.model_selection import train_test_split, GridSearchCV,
cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
import warnings # current version of seaborn generates a bunch of
warnings that we'll ignore
warnings.filterwarnings("ignore")

pd.read_csv("breast-cancer.csv")

```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean
area_mean \					
0	842302	M	17.99	10.38	122.80
1001.0					
1	842517	M	20.57	17.77	132.90
1326.0					
2	84300903	M	19.69	21.25	130.00
1203.0					
3	84348301	M	11.42	20.38	77.58
386.1					
4	84358402	M	20.29	14.34	135.10
1297.0					
..
...					
564	926424	M	21.56	22.39	142.00
1479.0					
565	926682	M	20.13	28.25	131.20
1261.0					
566	926954	M	16.60	28.08	108.30
858.1					
567	927241	M	20.60	29.33	140.10
1265.0					
568	92751	B	7.76	24.54	47.92
181.0					
	smoothness_mean	compactness_mean	concavity_mean	concave	
points_mean \					
0	0.11840	0.27760	0.30010		
0.14710					
1	0.08474	0.07864	0.08690		

0.07017			
2	0.10960	0.15990	0.19740
0.12790			
3	0.14250	0.28390	0.24140
0.10520			
4	0.10030	0.13280	0.19800
0.10430			
..
...			
564	0.11100	0.11590	0.24390
0.13890			
565	0.09780	0.10340	0.14400
0.09791			
566	0.08455	0.10230	0.09251
0.05302			
567	0.11780	0.27700	0.35140
0.15200			
568	0.05263	0.04362	0.00000
0.00000			

	...	radius_worst	texture_worst	perimeter_worst	area_worst	\
0	...	25.380	17.33	184.60	2019.0	
1	...	24.990	23.41	158.80	1956.0	
2	...	23.570	25.53	152.50	1709.0	
3	...	14.910	26.50	98.87	567.7	
4	...	22.540	16.67	152.20	1575.0	
..	
564	...	25.450	26.40	166.10	2027.0	
565	...	23.690	38.25	155.00	1731.0	
566	...	18.980	34.12	126.70	1124.0	
567	...	25.740	39.42	184.60	1821.0	
568	...	9.456	30.37	59.16	268.6	

	smoothness_worst	compactness_worst	concavity_worst	\
0	0.16220	0.66560	0.7119	
1	0.12380	0.18660	0.2416	
2	0.14440	0.42450	0.4504	
3	0.20980	0.86630	0.6869	
4	0.13740	0.20500	0.4000	
..	
564	0.14100	0.21130	0.4107	
565	0.11660	0.19220	0.3215	
566	0.11390	0.30940	0.3403	
567	0.16500	0.86810	0.9387	
568	0.08996	0.06444	0.0000	

	concave points_worst	symmetry_worst	fractal_dimension_worst
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758

3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678
...
564	0.2216	0.2060	0.07115
565	0.1628	0.2572	0.06637
566	0.1418	0.2218	0.07820
567	0.2650	0.4087	0.12400
568	0.0000	0.2871	0.07039

[569 rows x 32 columns]

```
df = pd.read_csv("breast-cancer.csv")
print("Dataset Shape:", df.shape)
print(df.head())
```

Dataset Shape: (569, 32)

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean \
0	842302	M	17.99	10.38	122.80	1001.0
1	842517	M	20.57	17.77	132.90	1326.0
2	84300903	M	19.69	21.25	130.00	1203.0
3	84348301	M	11.42	20.38	77.58	386.1
4	84358402	M	20.29	14.34	135.10	1297.0

	smoothness_mean	compactness_mean	concavity_mean	concave	points_mean \
0	0.11840	0.27760	0.3001		0.14710
1	0.08474	0.07864	0.0869		0.07017
2	0.10960	0.15990	0.1974		0.12790
3	0.14250	0.28390	0.2414		0.10520
4	0.10030	0.13280	0.1980		0.10430

	...	radius_worst	texture_worst	perimeter_worst	area_worst \
0	...	25.38	17.33	184.60	2019.0
1	...	24.99	23.41	158.80	1956.0
2	...	23.57	25.53	152.50	1709.0
3	...	14.91	26.50	98.87	567.7
4	...	22.54	16.67	152.20	1575.0

	smoothness_worst	compactness_worst	concavity_worst	concave
--	------------------	-------------------	-----------------	---------

points_worst \			
0	0.1622	0.6656	0.7119
0.2654			
1	0.1238	0.1866	0.2416
0.1860			
2	0.1444	0.4245	0.4504
0.2430			
3	0.2098	0.8663	0.6869
0.2575			
4	0.1374	0.2050	0.4000
0.1625			

	symmetry_worst	fractal_dimension_worst
0	0.4601	0.11890
1	0.2750	0.08902
2	0.3613	0.08758
3	0.6638	0.17300
4	0.2364	0.07678

[5 rows x 32 columns]

```
print("\nSummary Statistics:\n", df.describe())
```

Summary Statistics:

	id	radius_mean	texture_mean	perimeter_mean
area_mean \				
count	5.690000e+02	569.000000	569.000000	569.000000
569.000000				
mean	3.037183e+07	14.127292	19.289649	91.969033
654.889104				
std	1.250206e+08	3.524049	4.301036	24.298981
351.914129				
min	8.670000e+03	6.981000	9.710000	43.790000
143.500000				
25%	8.692180e+05	11.700000	16.170000	75.170000
420.300000				
50%	9.060240e+05	13.370000	18.840000	86.240000
551.100000				
75%	8.813129e+06	15.780000	21.800000	104.100000
782.700000				
max	9.113205e+08	28.110000	39.280000	188.500000
2501.000000				

	smoothness_mean	compactness_mean	concavity_mean	concave
points_mean \				
count	569.000000	569.000000	569.000000	
569.000000				
mean	0.096360	0.104341	0.088799	
0.048919				

std	0.014064	0.052813	0.079720
0.038803			
min	0.052630	0.019380	0.000000
0.000000			
25%	0.086370	0.064920	0.029560
0.020310			
50%	0.095870	0.092630	0.061540
0.033500			
75%	0.105300	0.130400	0.130700
0.074000			
max	0.163400	0.345400	0.426800
0.201200			

	symmetry_mean	...	radius_worst	texture_worst
perimeter_worst \				
count	569.000000	...	569.000000	569.000000
569.000000				
mean	0.181162	...	16.269190	25.677223
107.261213				
std	0.027414	...	4.833242	6.146258
33.602542				
min	0.106000	...	7.930000	12.020000
50.410000				
25%	0.161900	...	13.010000	21.080000
84.110000				
50%	0.179200	...	14.970000	25.410000
97.660000				
75%	0.195700	...	18.790000	29.720000
125.400000				
max	0.304000	...	36.040000	49.540000
251.200000				

	area_worst	smoothness_worst	compactness_worst
concavity_worst \			
count	569.000000	569.000000	569.000000
569.000000			
mean	880.583128	0.132369	0.254265
0.272188			
std	569.356993	0.022832	0.157336
0.208624			
min	185.200000	0.071170	0.027290
0.000000			
25%	515.300000	0.116600	0.147200
0.114500			
50%	686.500000	0.131300	0.211900
0.226700			
75%	1084.000000	0.146000	0.339100
0.382900			
max	4254.000000	0.222600	1.058000

1.252000

	concave points_worst	symmetry_worst	fractal_dimension_worst
count	569.000000	569.000000	569.000000
mean	0.114606	0.290076	0.083946
std	0.065732	0.061867	0.018061
min	0.000000	0.156500	0.055040
25%	0.064930	0.250400	0.071460
50%	0.099930	0.282200	0.080040
75%	0.161400	0.317900	0.092080
max	0.291000	0.663800	0.207500

[8 rows x 31 columns]

```
print("\nMissing Values:\n", df.isnull().sum())
```

Missing Values:

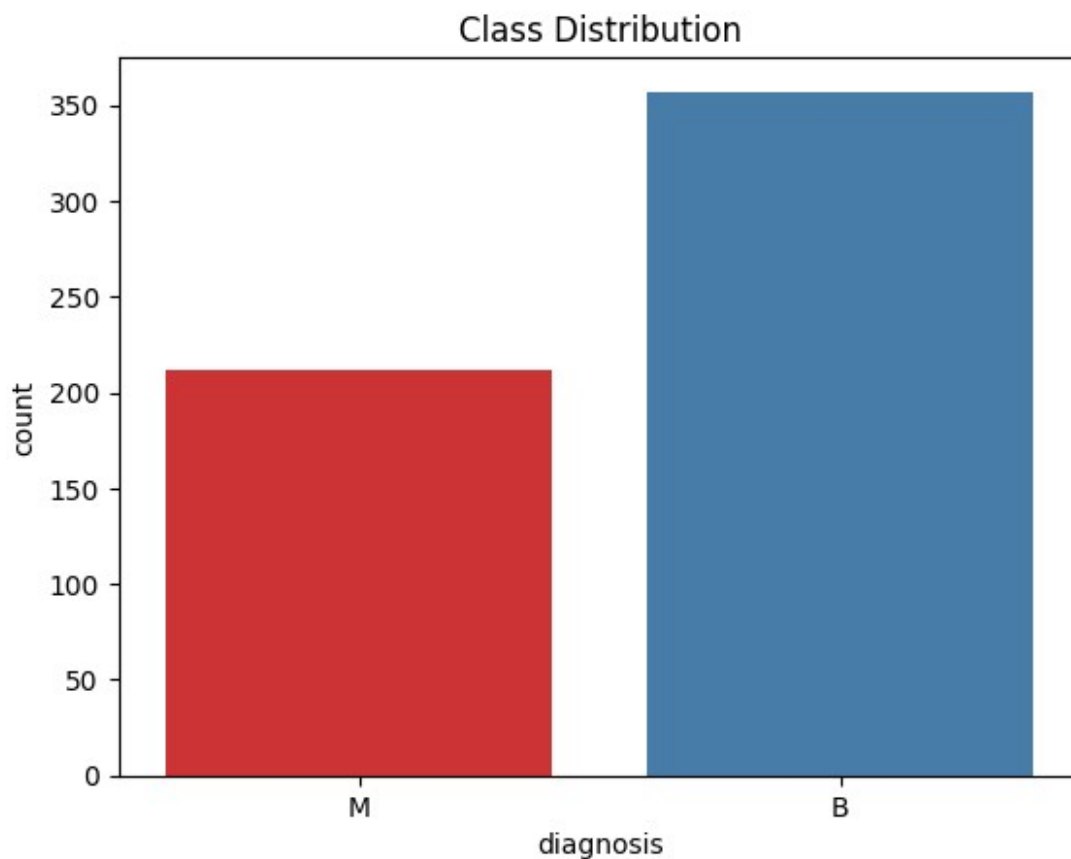
id	0
diagnosis	0
radius_mean	0
texture_mean	0
perimeter_mean	0
area_mean	0
smoothness_mean	0
compactness_mean	0
concavity_mean	0
concave points_mean	0
symmetry_mean	0
fractal_dimension_mean	0
radius_se	0
texture_se	0
perimeter_se	0
area_se	0
smoothness_se	0
compactness_se	0
concavity_se	0
concave points_se	0
symmetry_se	0
fractal_dimension_se	0
radius_worst	0
texture_worst	0
perimeter_worst	0
area_worst	0
smoothness_worst	0
compactness_worst	0
concavity_worst	0
concave points_worst	0
symmetry_worst	0

```
fractal_dimension_worst    0
dtype: int64
```

Visualization

```
sns.countplot(x='diagnosis', data=df, palette="Set1")
plt.title("Class Distribution")
plt.show()


print(df['diagnosis'].value_counts(normalize=True))
```



```
diagnosis
B    0.627417
M    0.372583
Name: proportion, dtype: float64



px.pie(df, 'diagnosis',
color='diagnosis',color_discrete_sequence=['#007500','#5CFF5C'],title=
'Data Distribution')
```




Data Distribution




■ B
■ M

```
for column in df.drop('diagnosis',axis=1).columns[:5]:  
    fig =  
    px.box(data_frame=df,x='diagnosis',color='diagnosis',y=column,color_discrete_sequence=['#007500','#5CFF5C'],orientation='v')  
    fig.show()
```






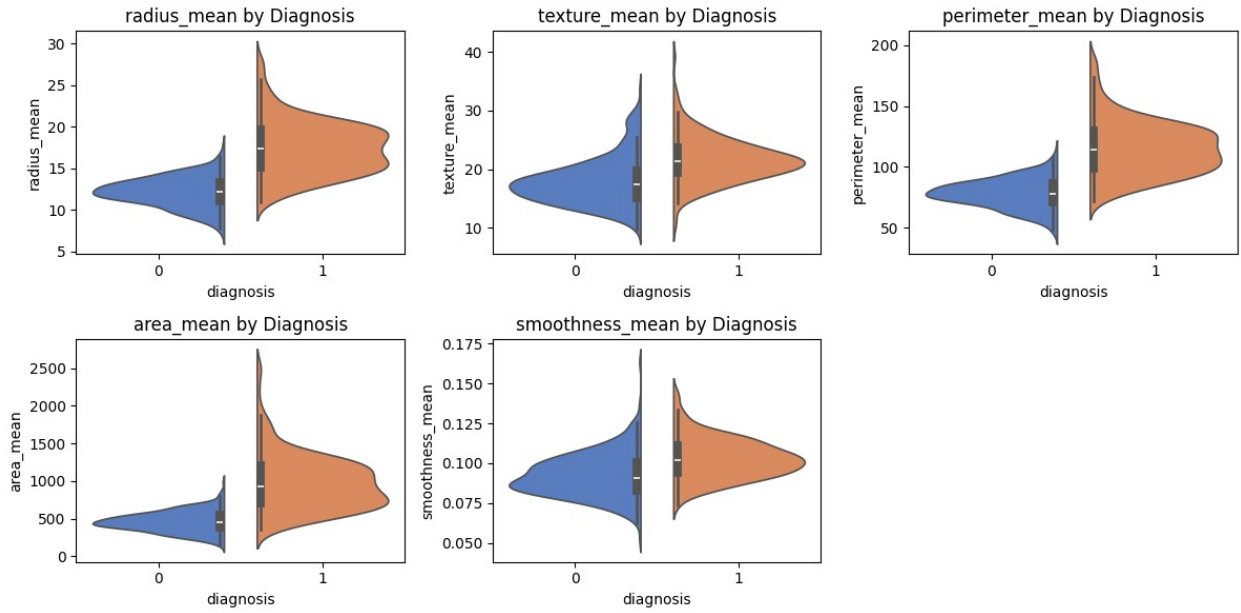
```
df['diagnosis'] = (df['diagnosis'] == 'M').astype(int) #encode the
label into 1/0
corr = df.corr()
plt.figure(figsize=(20,20))
sns.heatmap(corr, cmap='viridis_r',annot=True)
plt.show()
```




```
# Correlation with diagnosis
corr = df.corr()['diagnosis'].sort_values(ascending=False).head(11) #
top 10 + diagnosis
plt.figure(figsize=(8,6))
sns.heatmap(df[corr.index].corr(), annot=True, cmap="coolwarm",
fmt=".2f")
plt.title("Top Correlated Features with Diagnosis")
plt.show()
```



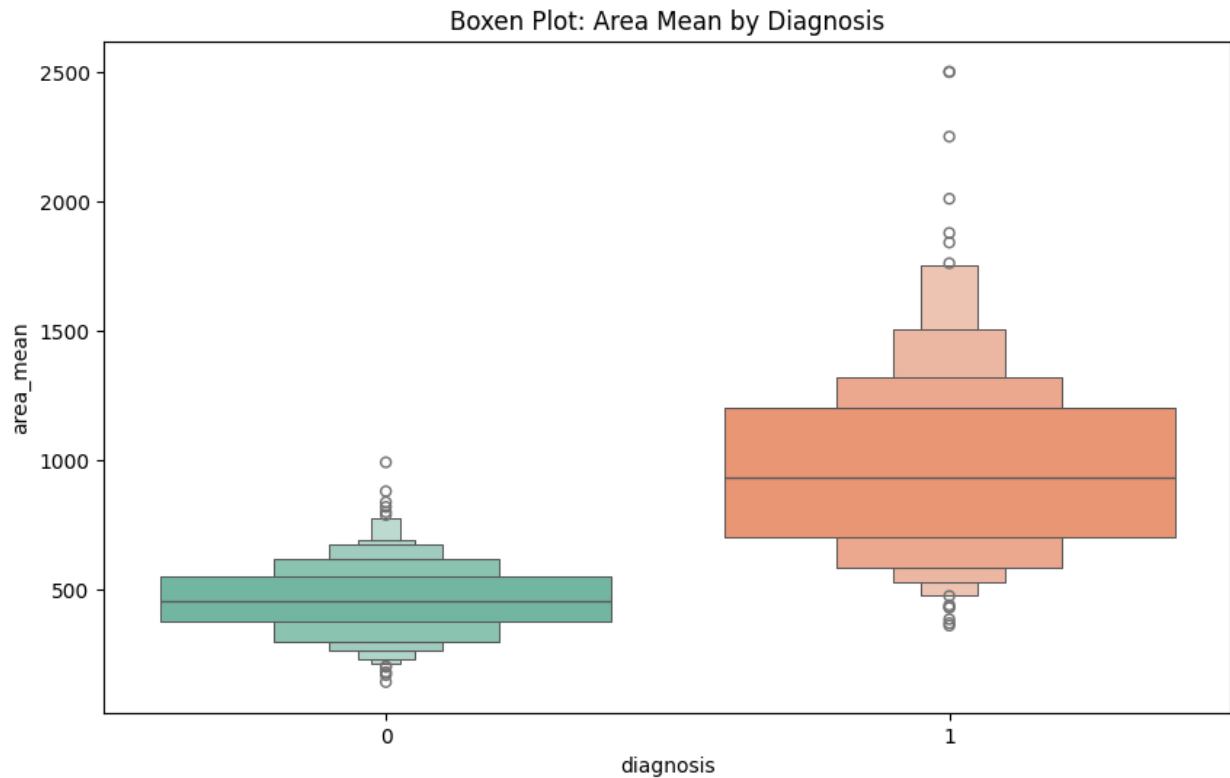
```
features = ['radius_mean', 'texture_mean', 'perimeter_mean',
            'area_mean', 'smoothness_mean']
plt.figure(figsize=(12,6))
for i, feat in enumerate(features, 1):
    plt.subplot(2, 3, i)
    sns.violinplot(x="diagnosis", y=feat, data=df, palette="muted",
split=True)
    plt.title(f"{feat} by Diagnosis")
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(10,6))
sns.swarmplot(x="diagnosis", y="radius_mean", data=df, palette="Set1",
alpha=0.7)
plt.title("Swarm Plot: Radius Mean by Diagnosis")
plt.show()
```

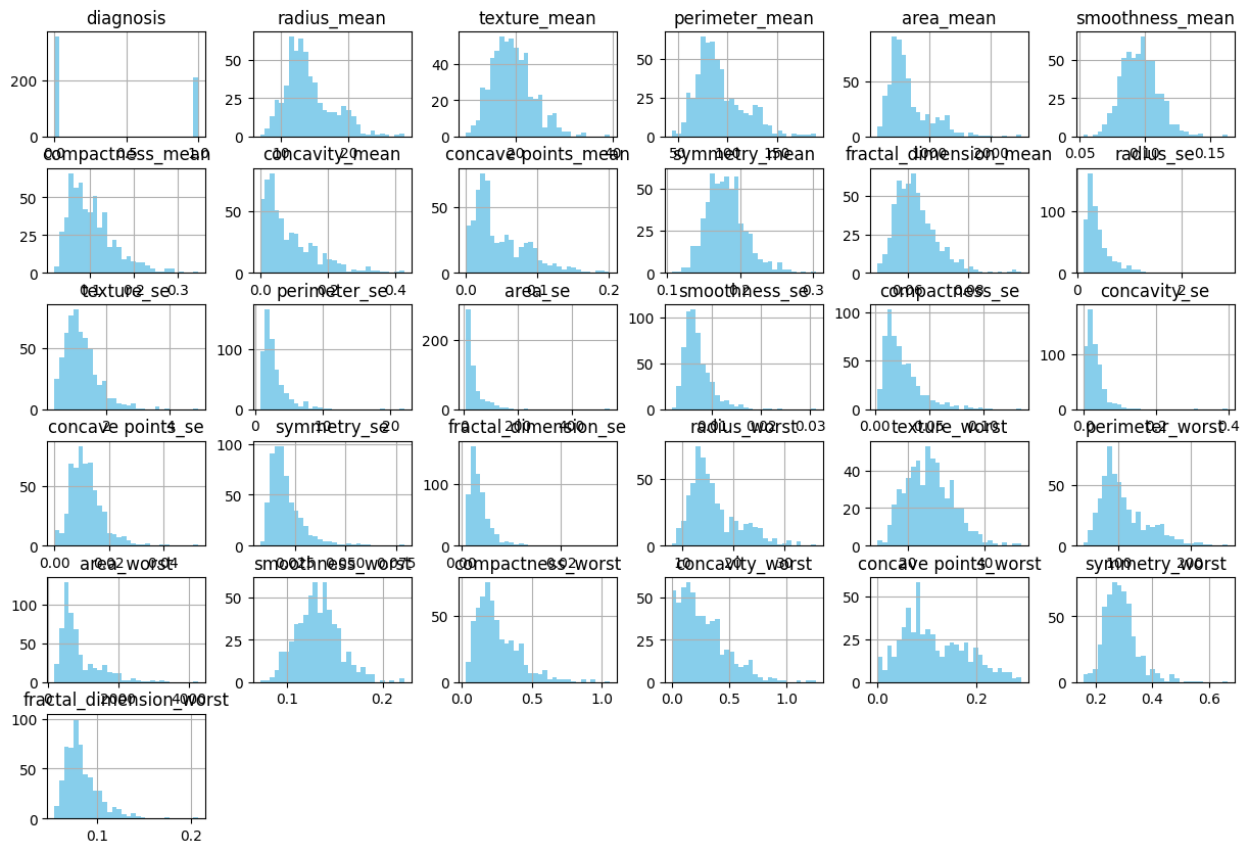


```
plt.figure(figsize=(10,6))
sns.boxenplot(x="diagnosis", y="area_mean", data=df, palette="Set2")
plt.title("Boxen Plot: Area Mean by Diagnosis")
plt.show()
```




```
df.drop(['id'], axis=1, errors='ignore').hist(bins=30,
figsize=(15,10), color="skyblue")
plt.suptitle("Feature Distributions", fontsize=16)
plt.show()
```

Feature Distributions




```
df['radius_bin'] = pd.qcut(df['radius_mean'], q=5) # divide into 5 bins
ct = pd.crosstab(df['radius_bin'], df['diagnosis'])

ct.plot(kind="bar", stacked=True, figsize=(10,6), colormap="coolwarm")
plt.title("Diagnosis Distribution across Radius Mean Bins")
plt.ylabel("Count")
plt.show()
```






```
features = ['radius_mean', 'texture_mean', 'perimeter_mean',
            'area_mean']
for feat in features:
    plt.figure(figsize=(6,4))
    sns.kdeplot(df[df['diagnosis']==0][feat], label="Benign",
                shade=True)
    sns.kdeplot(df[df['diagnosis']==1][feat], label="Malignant",
                shade=True)
    plt.title(f"Distribution of {feat} by Diagnosis")
    plt.legend()
    plt.show()
```


Distribution of radius_mean by Diagnosis




Distribution of texture_mean by Diagnosis



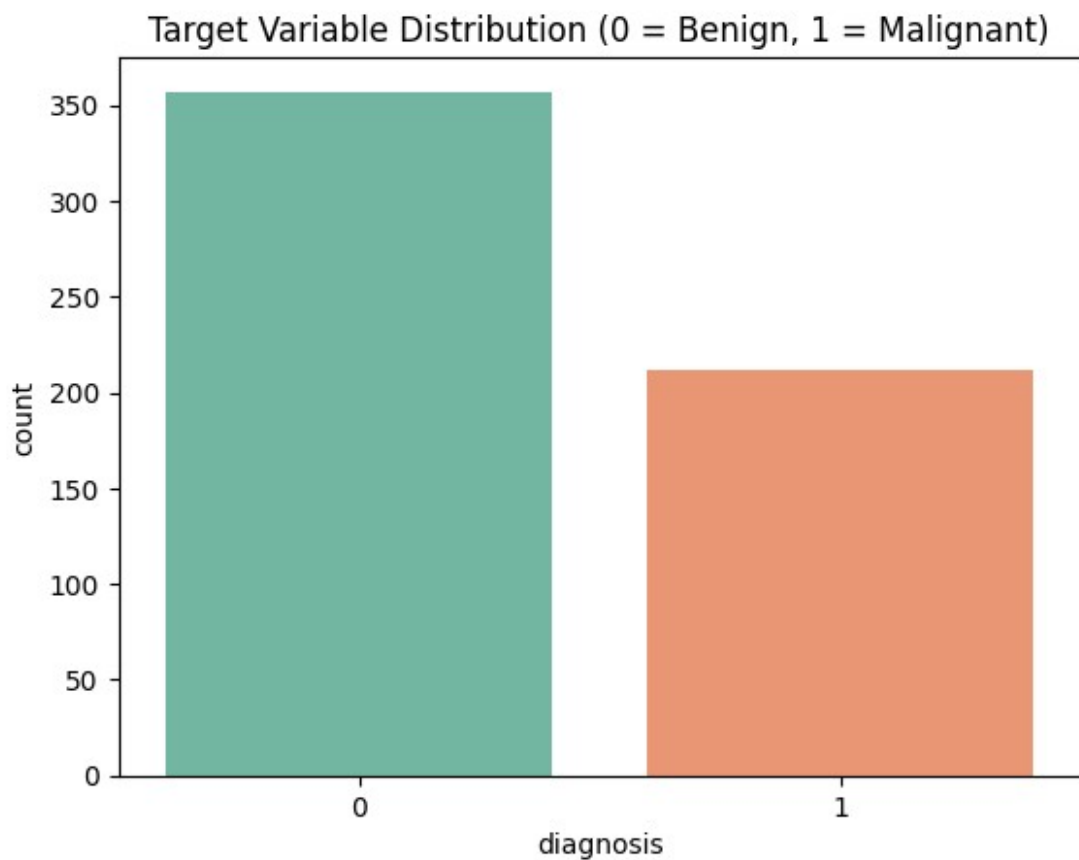


```
plt.figure(figsize=(10,6))  
sns.boxplot(x="diagnosis", y="radius_mean", data=df, palette="Set2")
```

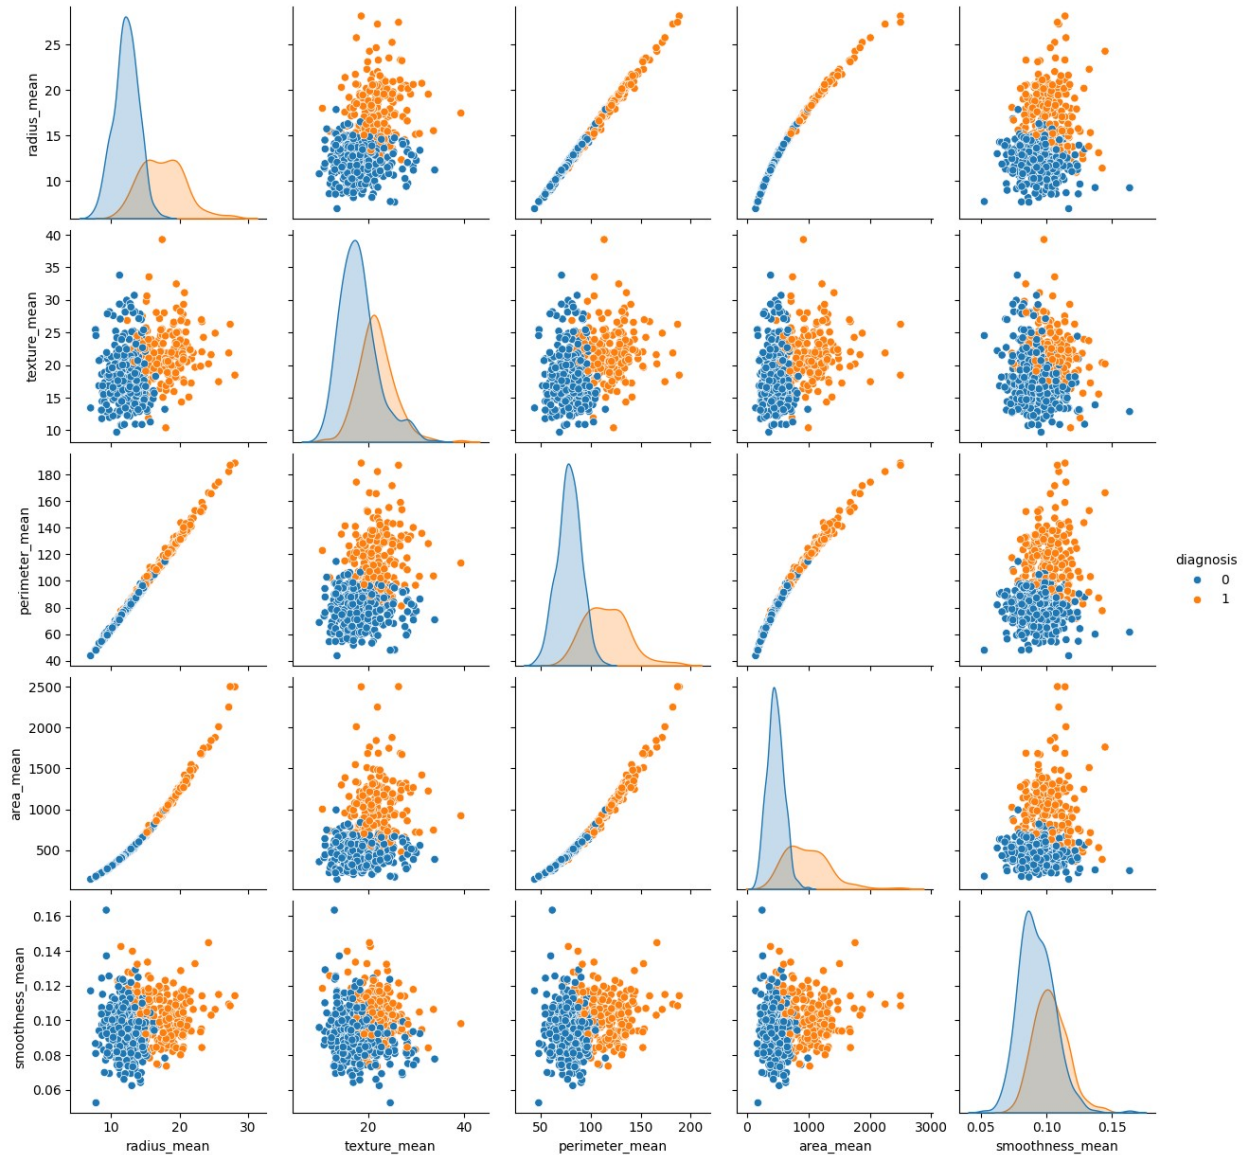
```
plt.title("Radius Mean by Diagnosis")
plt.show()
```



```
# Distribution of classes
sns.countplot(x='diagnosis', data=df, palette='Set2')
plt.title("Target Variable Distribution (0 = Benign, 1 = Malignant)")
plt.show()
```



```
# Select only first 5 numerical features + target column
sns.pairplot(df[['radius_mean', 'texture_mean', 'perimeter_mean',
'area_mean', 'smoothness_mean', 'diagnosis']],
             hue="diagnosis", diag_kind="kde")
plt.show()
```



```
# Preprocessing
```

```
X = df.drop(['id', 'diagnosis'], axis=1, errors='ignore') # drop id if
present
y = df['diagnosis']
```

```
X = X.select_dtypes(include=[np.number])
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)
```


```

from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0], X_pca[:,1], c=y, cmap="coolwarm",
            edgecolor="k", s=40)
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.title("PCA Visualization of Breast Cancer Data")
plt.show()

```



```

from mpl_toolkits.mplot3d import Axes3D
from sklearn.decomposition import PCA

pca = PCA(n_components=3)
X_pca = pca.fit_transform(X_scaled)

fig = plt.figure(figsize=(8,6))


```

```

ax = fig.add_subplot(111, projection='3d')
ax.scatter(X_pca[:,0], X_pca[:,1], X_pca[:,2], c=y, cmap="coolwarm",
s=50)
ax.set_xlabel("PC1")
ax.set_ylabel("PC2")
ax.set_zlabel("PC3")
plt.title("3D PCA Visualization")
plt.show()

```

3D PCA Visualization



```

# Train SVM (Linear & RBF)

svm_linear = SVC(kernel="linear")
svm_linear.fit(X_train, y_train)
y_pred_linear = svm_linear.predict(X_test)

print("\n--- Linear SVM Performance ---")
print(classification_report(y_test, y_pred_linear))

--- Linear SVM Performance ---

```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	72
1	1.00	0.90	0.95	42
accuracy			0.96	114
macro avg	0.97	0.95	0.96	114
weighted avg	0.97	0.96	0.96	114

```
svm_rbf = SVC(kernel="rbf")
svm_rbf.fit(X_train, y_train)
y_pred_rbf = svm_rbf.predict(X_test)

print("\n--- RBF SVM Performance ---")
print(classification_report(y_test, y_pred_rbf))
```

--- RBF SVM Performance ---				
	precision	recall	f1-score	support
0	0.96	1.00	0.98	72
1	1.00	0.93	0.96	42
accuracy			0.97	114
macro avg	0.98	0.96	0.97	114
weighted avg	0.97	0.97	0.97	114

Decision Boundary (2D Visualization)

Use only two features (e.g., radius_mean, texture_mean)

```
X2 = df[['radius_mean', 'texture_mean']]
```

```
y2 = df['diagnosis']
```

```
X2_scaled = scaler.fit_transform(X2)
```

```
X2_train, X2_test, y2_train, y2_test = train_test_split(X2_scaled, y2,
test_size=0.2, random_state=42)
```

```
clf2 = SVC(kernel='rbf', C=1, gamma=0.5)
```

```
clf2.fit(X2_train, y2_train)
```

```
SVC(C=1, gamma=0.5)
```

Meshgrid for decision boundary

```
x_min, x_max = X2_scaled[:,0].min() - 1, X2_scaled[:,0].max() + 1
```

```
y_min, y_max = X2_scaled[:,1].min() - 1, X2_scaled[:,1].max() + 1
```

```
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
                    np.linspace(y_min, y_max, 200))
```


```
Z = clf2.predict(np.c_[xx.ravel(), yy.ravel()])
```

```

Z = Z.reshape(xx.shape)

plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)
plt.scatter(X2_scaled[:,0], X2_scaled[:,1], c=y2, edgecolors='k',
            cmap=plt.cm.coolwarm)
plt.xlabel("Radius Mean")
plt.ylabel("Texture Mean")
plt.title("SVM Decision Boundary (RBF Kernel)")
plt.show()

```



```

# Hyperparameter Tuning

param_grid = {
    'C': [0.1, 1, 10, 100],
    'gamma': [1, 0.1, 0.01, 0.001],
    'kernel': ['rbf']
}

grid = GridSearchCV(SVC(), param_grid, refit=True, verbose=1, cv=5)
grid.fit(X_train, y_train)

print("\nBest Hyperparameters:", grid.best_params_)
print("Best Cross-Validation Score:", grid.best_score_)

```


Fitting 5 folds for each of 16 candidates, totalling 80 fits

Best Hyperparameters: {'C': 100, 'gamma': 0.01, 'kernel': 'rbf'}

Best Cross-Validation Score: 0.9758241758241759

Evaluate best model

```
y_pred_best = grid.best_estimator_.predict(X_test)
```

```
print("\n--- Tuned RBF SVM Performance ---")
```

```
print(classification_report(y_test, y_pred_best))
```

--- Tuned RBF SVM Performance ---

	precision	recall	f1-score	support
0	0.96	0.97	0.97	72
1	0.95	0.93	0.94	42
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

Cross-validation with Linear SVM

```
cv_scores = cross_val_score(svm_linear, X_scaled, y, cv=5)
```

```
print("\nCross-validation Accuracy (Linear SVM):", cv_scores.mean())
```

Cross-validation Accuracy (Linear SVM): 0.9701443875174661

```
from sklearn.feature_selection import SelectKBest, f_classif
```

```
selector = SelectKBest(score_func=f_classif, k=10)
```

```
fit = selector.fit(X_scaled, y)
```

```
feature_scores = pd.DataFrame({  
    'Feature': X.columns,  
    'Score': fit.scores_  
}).sort_values(by="Score", ascending=False)
```

```
print(feature_scores.head(10))
```

	Feature	Score
27	concave points_worst	964.385393
22	perimeter_worst	897.944219
7	concave points_mean	861.676020
20	radius_worst	860.781707
2	perimeter_mean	697.235272
23	area_worst	661.600206
0	radius_mean	646.981021
3	area_mean	573.060747

```
6         concavity_mean  533.793126
26        concavity_worst  436.691939
```

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay,
roc_curve, auc
```



```
# Confusion Matrix
```

```
cm = confusion_matrix(y_test, y_pred_best)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=["Benign", "Malignant"])
disp.plot(cmap="Blues")
plt.show()
```

```
# ROC Curve
```

```
fpr, tpr, _ = roc_curve(y_test, y_pred_best)
roc_auc = auc(fpr, tpr)
```

```
plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, label=f"ROC curve (area = {roc_auc:.2f})")
plt.plot([0,1],[0,1], '--', color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



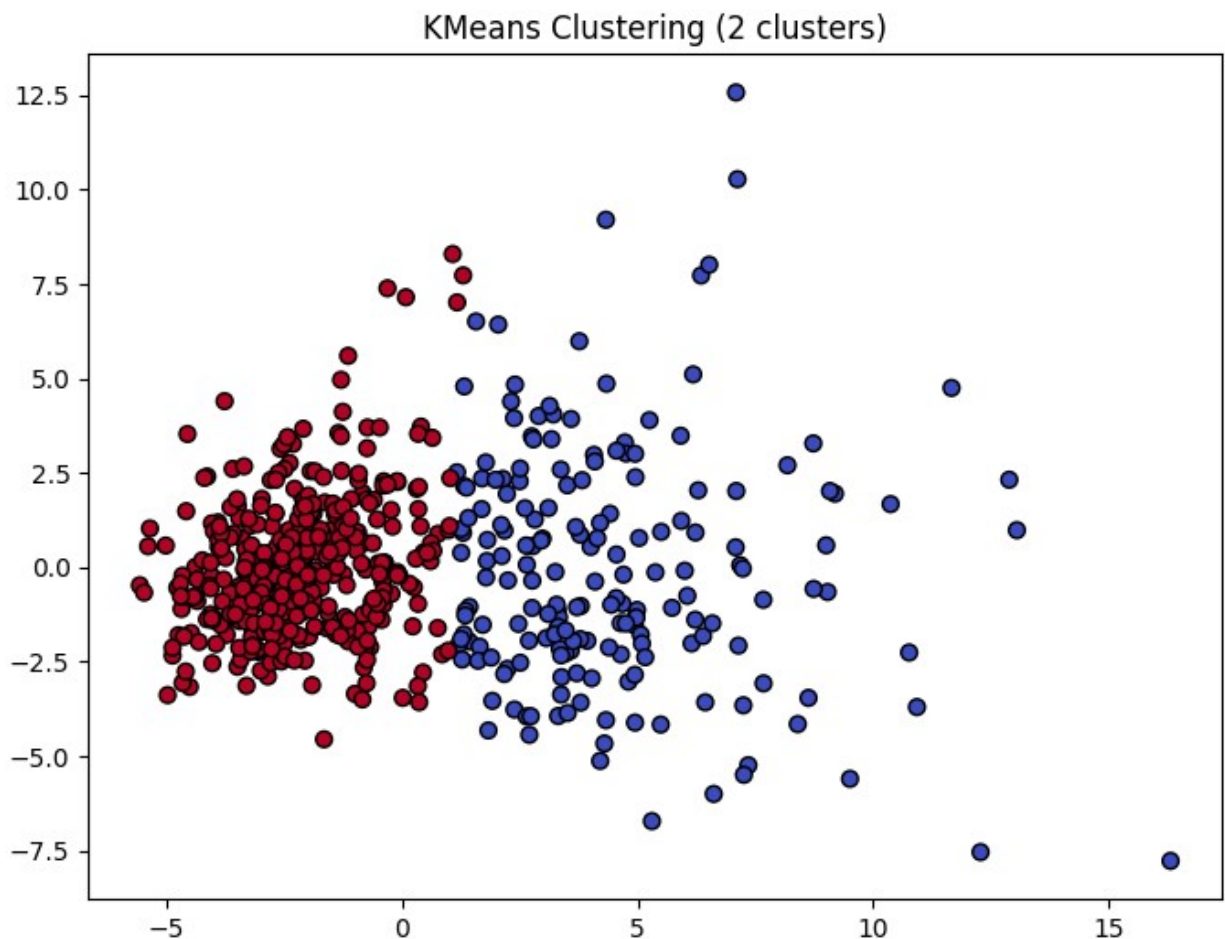
```

from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=2, random_state=42)
clusters = kmeans.fit_predict(X_scaled)

plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0], X_pca[:,1], c=clusters, cmap="coolwarm", s=40,
            edgecolor="k")
plt.title("KMeans Clustering (2 clusters)")
plt.show()

```



```

from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score

# =====
# Hyperparameter Tuning (RBF SVM)
# =====
param_grid = {

```

```

'C': [0.1, 1, 10, 100],          # Regularization parameter
'gamma': [1, 0.1, 0.01, 0.001], # Kernel coefficient
'kernel': ['rbf']
}

grid = GridSearchCV(SVC(probability=True), param_grid, refit=True,
verbose=1, cv=5, scoring='accuracy')
grid.fit(X_train, y_train)

print("□ Best Parameters:", grid.best_params_)
print("□ Best CV Score:", grid.best_score_)

Fitting 5 folds for each of 16 candidates, totalling 80 fits
□ Best Parameters: {'C': 100, 'gamma': 0.01, 'kernel': 'rbf'}
□ Best CV Score: 0.9758241758241759

# Evaluate on test set
y_pred_best = grid.best_estimator_.predict(X_test)

print("\n--- Tuned RBF SVM Performance ---")
print("Accuracy:", accuracy_score(y_test, y_pred_best))
print(classification_report(y_test, y_pred_best))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_best))

--- Tuned RBF SVM Performance ---
Accuracy: 0.956140350877193

```

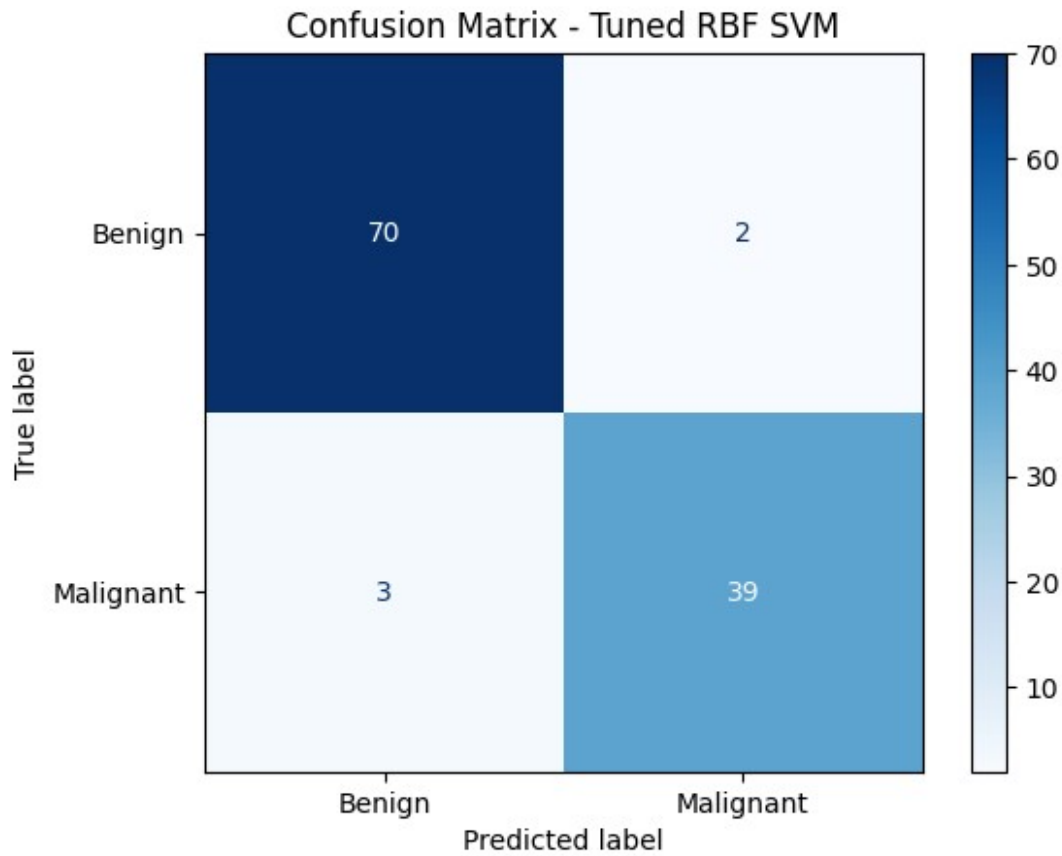
	precision	recall	f1-score	support
0	0.96	0.97	0.97	72
1	0.95	0.93	0.94	42
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

```

Confusion Matrix:
[[70  2]
 [ 3 39]]

# Plot Confusion Matrix
disp =
ConfusionMatrixDisplay(confusion_matrix=confusion_matrix(y_test,
y_pred_best),
                        display_labels=["Benign", "Malignant"])
disp.plot(cmap="Blues")
plt.title("Confusion Matrix - Tuned RBF SVM")
plt.show()

```



```
# Cross-validation (Linear SVM)
```

```
# Linear SVM
```

```
svm_linear = SVC(kernel="linear", C=1)
cv_scores_linear = cross_val_score(svm_linear, X_scaled, y, cv=5,
scoring='accuracy')
print("\nCross-validation Accuracy (Linear SVM):",
cv_scores_linear.mean())
```

```
Cross-validation Accuracy (Linear SVM): 0.9701443875174661
```

```
# Best RBF SVM (from GridSearch)
```

```
best_rbf = grid.best_estimator_
cv_scores_rbf = cross_val_score(best_rbf, X_scaled, y, cv=5,
scoring='accuracy')
print("Cross-validation Accuracy (Best RBF SVM):",
cv_scores_rbf.mean())
```

```
Cross-validation Accuracy (Best RBF SVM): 0.9683744760130415
```

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

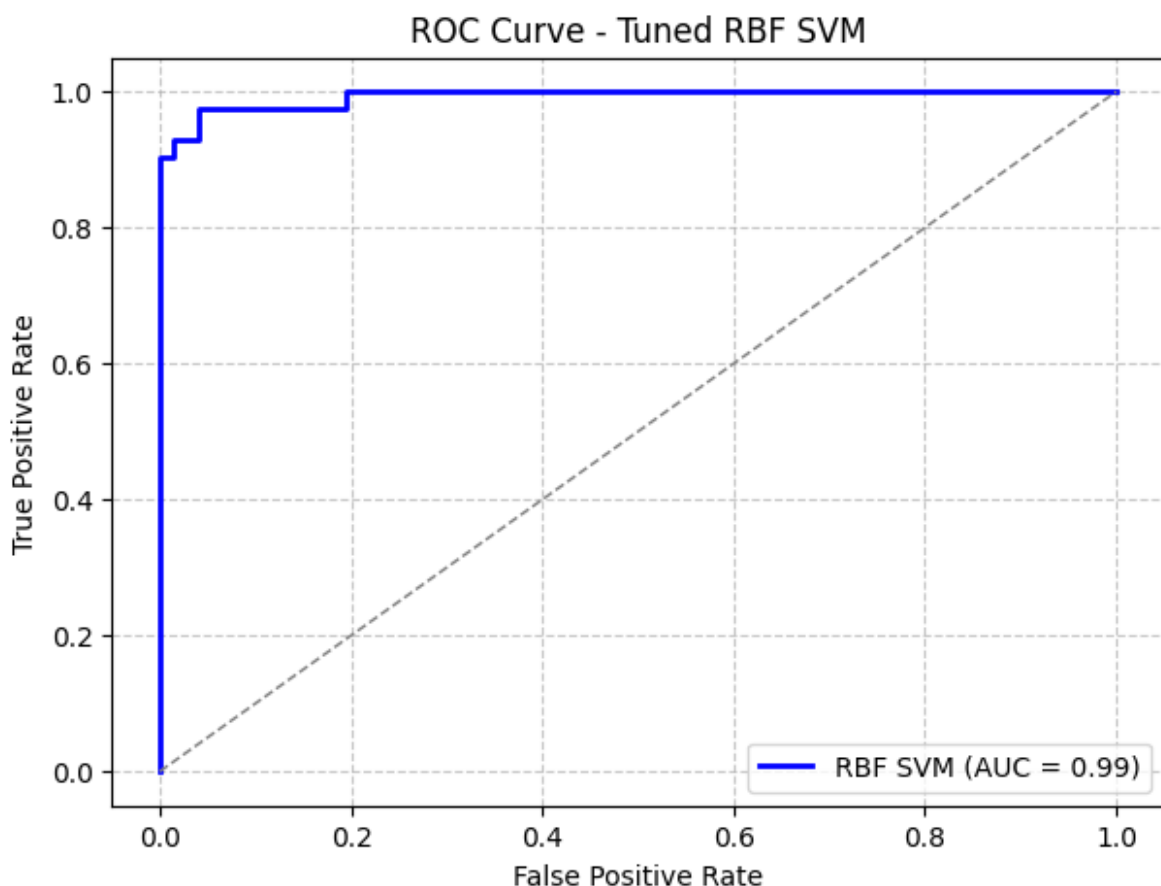
```

# Predict probabilities with best RBF SVM
y_prob = grid.best_estimator_.predict_proba(X_test)[: ,1]

# Compute ROC curve and AUC
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

# Plot ROC Curve
plt.figure(figsize=(7,5))
plt.plot(fpr, tpr, color="blue", lw=2, label=f"RBF SVM (AUC = {roc_auc:.2f})")
plt.plot([0,1], [0,1], color="gray", linestyle="--", lw=1)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Tuned RBF SVM")
plt.legend(loc="lower right")
plt.grid(True, linestyle="--", alpha=0.7)
plt.show()

```



```

# Cross-validation accuracy results
linear_acc = cv_scores_linear.mean()

```

```

rbf_acc = cv_scores_rbf.mean()

# Bar chart
models = ['Linear SVM', 'Tuned RBF SVM']
scores = [linear_acc, rbf_acc]

plt.figure(figsize=(6,5))
plt.bar(models, scores, color=['skyblue', 'salmon'])
plt.ylabel("Mean CV Accuracy")
plt.ylim(0,1)
for i, score in enumerate(scores):
    plt.text(i, score+0.01, f"{score:.3f}", ha='center', fontsize=11,
             fontweight='bold')
plt.title("Cross-Validation Accuracy Comparison")
plt.show()

```

